

Chapter 9: Virtual Memory

- Background
- Demand Paging
- Copy-on-Write
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- Operating-System Examples

Background

- **Virtual memory** – separation of user logical memory from physical memory.
 - ☞ Only part of the program needs to be in memory for execution.
 - ☞ Logical address space can therefore be much larger than physical address space.

- ☞ Allows address spaces to be shared by several processes.


- ☞ Allows for more efficient process creation.


- Virtual memory can be implemented via:


- ☞ Demand paging

- ☞ Demand segmentation

Virtual Memory That is Larger Than Physical Memory








Demand Paging

■ Bring a page into memory only when it is needed.

- ➡ Less I/O needed
- ➡ Less memory needed
- ➡ Faster response
- ➡ More users

- Page is needed $\Rightarrow$ reference to it
 - ☞ invalid reference $\Rightarrow$ abort
 - ☞ not-in-memory $\Rightarrow$ bring to memory

Transfer of a Paged Memory to Contiguous Disk Space



Valid-Invalid Bit

- With each page table entry a valid–invalid bit is associated
(1 $\Rightarrow$ in-memory, 0 $\Rightarrow$ not-in-memory)
- Initially valid–invalid, is set to 0 on all entries.
- Example of a page table snapshot.
- During address translation, if valid – invalid bit in page table entry is 0 $\Rightarrow$ page fault.

Valid-Invalid Bit

Frame # valid-invalid bit



	1
	1
	1
	1
	0
⋮	
	0
	0

page table

Page Table When Some Pages Are Not in Main Memory

0	A
1	B
2	C
3	D
4	E
5	F
6	G
7	H


logical
memory

	frame	valid-invalid bit
0	4	v
1		i
2	6	v
3		i
4		i
5	9	v
6		i
7		i

page table

0	
1	
2	
3	
4	A
5	
6	C
7	
8	
9	F
10	
11	
12	
13	
14	
15	

physical memory




Page Fault

- If there is ever a reference to a page, first reference will trap to OS $\Rightarrow$ page fault
- OS looks at another table to decide:
 - ☞ Invalid reference $\Rightarrow$ abort.
 - ☞ Just not in memory.
- Get empty frame.
- Swap page into frame.


- Reset tables, validation bit = 1.
- Restart instruction: Least Recently Used

👉 block move



👉 auto increment/decrement location

Steps in Handling a Page Fault



Performance of Demand Paging

■ Page Fault Rate $0 \leq p \leq 1.0$

☞ if $p = 0$ no page faults

☞ if $p = 1$, every reference is a fault

■ Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} = & (1 - p) \times \text{memory access} \\ & + p \text{ (page fault overhead} \\ & + [\text{swap page out}] \\ & + \text{swap page in} \\ & + \text{restart overhead}) \end{aligned}$$

Demand Paging Example

- Memory access time = 1 microsecond
- 50% of the time the page that is being replaced has been modified and therefore needs to be swapped out.
- Swap Page Time = 10 msec = 10,000 msec

$$\text{EAT} = \frac{(1 - p) \times 1 + p (15000)}{1 + 15000P} \quad (\text{in msec})$$

Process Creation

■ Virtual memory allows other benefits during process creation:

- Copy-on-Write
- Memory-Mapped Files

Copy-on-Write

- Copy-on-Write (COW) allows both parent and child processes to initially *share* the same pages in memory.

If either process modifies a shared page, only then is the page copied.

- COW allows more efficient process creation as only modified pages are copied.
- Free pages are allocated from a *pool* of zeroed-out pages.

What happens if there is no free frame?


- Page replacement – find some page in memory, but not really in use, swap it out.
 - ☞ algorithm
 - ☞ performance – want an algorithm which will result in minimum number of page faults.
- Same page may be brought into memory several times.

Page Replacement

- Prevent over-allocation of memory by modifying page-fault service routine to include page replacement.
- Use *modify (dirty) bit* to reduce overhead of page transfers – only modified pages are written to disk.

- Page replacement completes separation between logical memory and physical memory – large virtual memory can be provided on a smaller physical memory.

Need For Page Replacement





Basic Page Replacement

1. Find the location of the desired page on disk.
2. Find a free frame:
 - If there is a free frame, use it.
 - If there is no free frame, use a page replacement algorithm to select a *victim* frame.

- Read the desired page into the (newly) free frame. Update the page and frame tables.
- Restart the process.


Page Replacement



Page Replacement Algorithms

- Want lowest page-fault rate.
- Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string.
- In all our examples, the reference string is 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5.

Graph of Page Faults Versus The Number of Frames



First-In-First-Out (FIFO) Algorithm

先进先出页面置换算法

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

- 3 frames (3 pages can be in memory at a time per process)

1	1	4	5
2	2	1	3
3	3	2	4

9 page faults

4 frames Reference string:

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1	1	5	4
2	2	1	5
3	3	2	
4	4	3	

10 page faults

FIFO Replacement – Belady’s Anomaly
(Belady異常)

more frames $\Rightarrow$ less page faults


FIFO Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2	2	2	4	4	4	0	0	0	7	7	7
	0	0	0	3	3	3	2	2	2	1	1	1	0	0
		1	1	1	0	0	0	3	3	3	2	2	2	1

FIFO Illustrating Belady's Anomaly




Optimal Algorithm

最优页面置换算法

- Replace page that will not be used for longest period of time.
- 4 frames example

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



6 page faults

- How do you know this?
- Used for measuring how well your algorithm performs.

Optimal Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2	2		2		2							7		
	0	0	0		0	4		0		0							0		
		1	1		3	3		3		1							1		


Least Recently Used (LRU) Algorithm

最近最少使用页面置换算法

■ Reference string:

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

▲ ▲ ▲ ▲ ✓ ✓ ▲ ✓ ✓ ▲ ▲ ▲



Reference string: 1, 2, 3, 4, 1, 2, **5**, 1, 2, **3**, **4**, **5**

1	1	1	1	5
2	2	2	2	2
3	5	5	4	4
4	4	3	3	3

8 page faults

■ Counter implementation

- ➡ Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter.
- ➡ When a page needs to be changed, look at the counters to determine which are to change.

LRU Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

LRU Algorithm (Cont.)

■ Stack implementation – keep a stack of page numbers in a double link form:

☞ Page referenced:

 move it to the top


 requires 6 pointers to be changed

☞ No search for replacement

Use Of A Stack to Record The Most Recent Page References

reference string


4 7 0 7 1 0 1 2 1 2 7 1 2



stack before a



stack after b



LRU Approximation Algorithms


■ Reference bit

- ☞ With each page associate a bit, initially = 0
- ☞ When page is referenced bit set to 1.
- ☞ Replace the one which is 0 (if one exists). We do not know the order, however.


■ Second chance

- ☞ Need reference bit.
- ☞ Clock replacement.
- ☞ If page to be replaced (in clock order) has reference bit = 1. then:
 - 📄 set reference bit 0.
 - 📄 leave page in memory.
 - 📄 replace next page (in clock order), subject to same rules.

Second-Chance (clock) Page-Replacement Algorithm



(a)



(b)

Counting Algorithms

- Keep a counter of the number of references that have been made to each page.
- LFU Algorithm: replaces page with smallest count.
- MFU Algorithm: based on the argument that the page with the smallest count was probably just brought in and has yet to be used.

Allocation of Frames

- Each process needs **minimum** number of pages.
- Example: IBM – 6 pages to handle SS MOVE instruction:
 - ☞ instruction is 6 bytes, might span 2 pages.
 - ☞ 2 pages to handle **from**.
 - ☞ 2 pages to handle **to**.

■ Two major allocation schemes.

- ➡ fixed allocation

- ➡ priority allocation

Fixed Allocation

- Equal allocation – e.g., if 100 frames and 5 processes, give each 20 pages.
- Proportional allocation – Allocate according to the size of process.

s_i = size of process p_i

$$S = \sum s_i$$

m = total number of frames

$$a_i = \text{allocation for } p_i = \frac{s_i}{S} \times m$$

$$m = 64$$

$$s_i = 10$$

$$s_2 = 127$$

$$a_1 = \frac{10}{137} \times 64 \approx 5$$

$$a_2 = \frac{127}{137} \times 64 \approx 59$$

Priority Allocation

- Use a proportional allocation scheme using priorities rather than size.
- If process P_i generates a page fault,
 - ☞ select for replacement one of its frames.
 - ☞ select for replacement a frame from a process with lower priority number.

Global vs. Local Allocation


- **Global** replacement – process selects a replacement frame from the set of all frames; one process can take a frame from another.
- **Local** replacement – each process selects from only its own set of allocated frames.

Thrashing

- If a process does not have “enough” pages, the page-fault rate is very high. This leads to:
 - ☞ low CPU utilization.
 - ☞ operating system thinks that it needs to increase the degree of multiprogramming.
 - ☞ another process added to the system.

■ **Thrashing** $\equiv$ a process is busy swapping pages in and out.

Thrashing



■ Why does paging work?

Locality model


☞ Process migrates from one locality to another.

☞ Localities may overlap.

■ Why does thrashing occur?

Σ size of locality > total memory size

Locality In A Memory-Reference Pattern



Working-Set Model

- $\Delta \equiv$ working-set window $\equiv$ a fixed number of page references

Example: 10,000 instruction

- WSS_i (working set of Process P_i) = total number of pages referenced in the most recent Δ (varies in time)

☞ if Δ too small will not encompass entire locality.

☞ if Δ too large will encompass several localities.

☞ if $\Delta = \infty \Rightarrow$ will encompass entire program.

■ $D = \sum WSS_i \equiv$ total demand frames


■ if $D > m \Rightarrow$ Thrashing

■ Policy if $D > m$, then suspend one of the processes.


Working-set model

page reference table

... 2 6 1 5 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...



$$WS(t_1) = \{1, 2, 5, 6, 7\}$$



$$WS(t_2) = \{3, 4\}$$


Keeping Track of the Working Set

- Approximate with interval timer + a reference bit
- Example: $\Delta = 10,000$
 - ☞ Timer interrupts after every 5000 time units.
 - ☞ Keep in memory 2 bits for each page.

- ☞ Whenever a timer interrupts copy and sets the values of all reference bits to 0.
- ☞ If one of the bits in memory = 1 $\Rightarrow$ page in working set.

- Why is this not completely accurate?
- Improvement = 10 bits and interrupt every 1000 time units.

Page-Fault Frequency Scheme



■ Establish “acceptable” page-fault rate.


- ☞ If actual rate too low, process loses frame.
- ☞ If actual rate too high, process gains frame.

Memory-Mapped Files


- Memory-mapped file I/O allows file I/O to be treated as routine memory access by *mapping* a disk block to a page in memory.
- A file is initially read using demand paging. A page-sized portion of the file is read from the file system into a physical page. Subsequent reads/writes to/from the file are treated as ordinary memory accesses.

- Simplifies file access by treating file I/O through memory rather than **read() write()** system calls.
- Also allows several processes to map the same file allowing the pages in memory to be shared.

Memory Mapped Files



Memory-Mapped Shared Memory in Windows



Allocating Kernel Memory


- Treated differently from user memory
- Often allocated from a free-memory pool
 - ☞ Kernel requests memory for structures of varying sizes
 - ☞ Some kernel memory needs to be contiguous

Buddy System


- Allocates memory from fixed-size segment consisting of physically-contiguous pages
- Memory allocated using **power-of-2 allocator**

- ☞ Satisfies requests in units sized as power of 2
- ☞ Request rounded up to next highest power of 2
- ☞ When smaller allocation needed than is available, current chunk split into two buddies of next-lower power of 2
 - 📄 Continue until appropriate sized chunk available

Buddy System Allocator



Slab Allocation



Other Considerations

■ Prepaging

- ☞ To reduce the large number of page faults that occurs at process startup
- ☞ Prepage all or some of the pages a process will need, before they are referenced
- ☞ But if prepaged pages are unused, I/O and memory was wasted

➡ Assume s pages are prepaged and α of the pages is used

📄 Is cost of $s * \alpha$ save pages faults $>$ or $<$ than the cost of prepaging $s * (1 - \alpha)$ unnecessary pages?

📄 α near zero $\Rightarrow$ prepaging loses

■ Page size selection

☞ fragmentation

☞ table size

☞ I/O overhead

☞ locality

Other Considerations (Cont.)

- **TLB Reach** - The amount of memory accessible from the TLB.
- $\text{TLB Reach} = (\text{TLB Size}) \times (\text{Page Size})$
- Ideally, the working set of each process is stored in the TLB. Otherwise there is a high degree of page faults.

Increasing the Size of the TLB

- **Increase the Page Size.** This may lead to an increase in fragmentation as not all applications require a large page size.
- **Provide Multiple Page Sizes.** This allows applications that require larger page sizes the opportunity to use them without an increase in fragmentation.

Other Considerations (Cont.)

Program structure

➡ **int A[][] = new int[1024][1024];**

➡ Each row is stored in one page

➡ Program 1 **for (j = 0; j < A.length; j++)
 for (i = 0; i < A.length; i++)
 A[i,j] = 0;**


1024 x 1024 page faults

➡ Program 2 **for (i = 0; i < A.length; i++)
 for (j = 0; j < A.length; j++)
 A[i,j] = 0;**

Other Considerations (Cont.)

- **I/O Interlock** – Pages must sometimes be locked into memory.
- Consider I/O. Pages that are used for copying a file from a device must be locked from being selected for eviction by a page replacement algorithm.

Reason Why Frames Used For I/O Must Be In Memory



Operating System Examples

- Windows

- Solaris

Windows

- Uses demand paging with **clustering**.
Clustering brings in pages surrounding the faulting page.
- Processes are assigned **working set minimum** and **working set maximum**
- Working set minimum is the minimum number of pages the process is guaranteed to have in memory


- A process may be assigned as many pages up to its working set maximum
- When the amount of free memory in the system falls below a threshold, **automatic working set trimming** is performed to restore the amount of free memory
- Working set trimming removes pages from processes that have pages in excess of their working set minimum

Solaris

- Maintains a list of free pages to assign faulting processes
- *Lotsfree* – threshold parameter (amount of free memory) to begin paging
- *Desfree* – threshold parameter to increasing paging
- *Minfree* – threshold parameter to being swapping

- Paging is performed by *pageout* process
- Pageout scans pages using modified clock algorithm
- *Scanrate* is the rate at which pages are scanned. This ranges from *slowscan* to *fastscan*
- Pageout is called more frequently depending upon the amount of free memory available

Solaris Page Scanner



- 某计算机主存按字节编址，逻辑地址和物理地址都是**32** 位，页表项大小为**4**字节。请回答下列问题。
- 若使用一级页表的分页存储管理方式，逻辑地址结构为：
- 页号（**20**位） 页内偏移量（**12**位）
- 则该页的大小是多少字节？
- 该页表最大占用多少字节？

- (1) 因为页内偏移量是12位，所以页大小为4 KB。
- (2) 页表项数为 $2^{32}/4K=2^{20}$ ，该一级页表最大为 $2^{20} \times 4 \text{ B}=4 \text{ MB}$ 。

■ 分区分配内存管理方式的主要保护措施是

- A. 界地址保护
- B. 程序代码保护
- C. 数据保护
- D. 栈保护

■ A。

考查分区分配存储管理方式的保护措施。分区分配存储管理方式的保护措施是设置界地址寄存器。每个进程都有自己独立的进程空间，如果一个进程在运行时所产生的地址在其地址空间之外，则发生地址越界，即当程序要访问某个内存单元时，由硬件检查是否允许，如果允许则执行，否则产生地址越界中断，由操作系统进行相应处理。

■ 一个分段存储管理系统中，地址长度为32位，其中段号占8位，则最大段长是：

- A. 2^8B
- B. 2^{16}B
- C. 2^{24}B
- D. 2^{32}B

■ C。

考查分段存储管理系统，分页式存储系统也一样。段地址为**32**位二进制数，其中**8**位表示段号，则段内位移占用 $32 - 8 = 24$ 位二进制数，故最大段长为 2^{24}B 。

■ 某基于动态分区存储管理的计算机，其主存容量为**55MB**（初始为空闲），采用最佳适配（**Best Fit**）算法，分配和释放的顺序为：分配**15MB**，分配**30MB**，释放**15MB**，分配**8MB**，分配**6MB**，此时主存中最大空闲分区的大小是。

- **A. 7MB**
- **B. 9MB**
- **C. 10MB**
- **D. 15MB**

- **B。考查动态分区分配。**考生需对动态分区分配的四种算法加以理解。首次适配法(**first-fit**): 按分区在内存的先后次序从头查找, 找到符合要求的第一个分区进行分配。
- **最佳适配法(best-fit):** 按分区在内存的先后次序从头查找, 找到其大小与要求相差最小的空闲分区进行分配。**最坏适配法(worst-fit):** 按分区在内存的先后次序从头查找, 找到最大的空闲分区进行分配。

- Assume a page reference string for a process with m frames (initially all empty). The page reference string has length p with n distinct page numbers occurring in it. For any page-replacement algorithms,
 - a. What is a lower bound on the number of page faults?
 - b. What is an upper bound on the number of page faults?

■ Answer:

■ a. n

■ b. p

- A certain computer provides its users with a virtual-memory space of 2^{32} bytes. The computer has 2^{18} bytes of physical memory. The virtual memory is implemented by paging, and the page size is 4096 bytes. A user process generates the virtual address.
- The virtual address in binary form is
- 0001 0001 0001 0010 0011 0100 0101 0110
- page size, the page table size, page number, offset

■ 页号 页内偏移量

■ **page size, the page table size, page number, offset**

■ **Since the page size is 2^{12} , the page table size is 2^{20} .**

■ **Offset, Therefore the low-order 12 bits
“0100 0101 0110”**

■ **page number , while the remaining 20 bits**

“0001 0001 0001 0010 0011”

- In a demand paging system, the page size is 1k bytes, and the page table is as follows(assuming use decimal values), translate a logical address 10 into its corresponding physical address, and why?

frame ↵
2 ↵
1 ↵
6 ↵

■ A. 8202

■ B. 4106

■ C. 2058

■ D. 1034

Chapter 10: File-System Interface

- File Concept
- Access Methods
- Directory Structure
- File System Mounting (文件系统安装)
- File Sharing (文件共享)
- Protection

File Concept

- Contiguous logical address space
(连续逻辑地址空间)

- Types:

 - ☞ Data

 - 📄 numeric

 - 📄 character

 - 📄 binary

 - ☞ Program

File Structure

- None - sequence of words, bytes
- Simple record structure
 - ➡ Lines
 - ➡ Fixed length
 - ➡ Variable length

■ Complex Structures

- ☞ Formatted document

- ☞ Relocatable load file

■ Can simulate last two with first method by inserting appropriate control characters.

■ Who decides:

- ☞ Operating system

- ☞ Program

File Attributes

- **Name** – only information kept in human-readable form.
- **Type** – needed for systems that support different types.
- **Location** – pointer to file location on device.
- **Size** – current file size.

- **Protection** – controls who can do reading, writing, executing.
- **Time, date, and user identification** – data for protection, security, and usage monitoring.
- Information about files are kept in the directory structure, which is maintained on the disk.

File Operations

- Create
- Write
- Read
- Reposition (重新定位) within file – file seek

- Delete
- Truncate
- $\text{Open}(F_i)$ – search the directory structure on disk for entry F_i , and move the content of entry to memory.
- $\text{Close}(F_i)$ – move the content of entry F_i in memory to directory structure on disk.

File Types – Name, Extension

file type	usual extension	function
executable	exe, com, bin or none	read to run machine- language program
object	obj, o	compiled, machine language, not linked
source code	c, cc, java, pas, asm, a	source code in various languages
batch	bat, sh	commands to the command interpreter
text	txt, doc	textual data, documents
word processor	wp, tex, rrf, doc	various word-processor formats
library	lib, a, so, dll, mpeg, mov, rm	libraries of routines for programmers
print or view	arc, zip, tar	ASCII or binary file in a format for printing or viewing
archive	arc, zip, tar	related files grouped into one file, sometimes com- pressed, for archiving or storage
multimedia	mpeg, mov, rm	binary file containing audio or A/V information

Access Methods

■ Sequential Access（顺序访问）

read next

write next

reset

no read after last write
(rewrite)

■ Direct Access（直接访问）

read n

write n

position to n


read next

write next

rewrite n

n = relative block number

Sequential-access File



Simulation of Sequential Access on a Direct-access File

sequential access	implementation for direct access
<i>reset</i>	<i>cp = 0;</i>
<i>read next</i>	<i>read cp;</i> <i>cp = cp+1;</i>
<i>write next</i>	<i>write cp;</i> <i>cp = cp+1;</i>

Example of Index and Relative Files

last name	logical record number
Adams	
Arthur	
Asher	
⋮	
Smith	

index file

Smith, John	social-security	age


relative file

Directory Structure

- A collection of nodes containing information about all files.


Directory

Files



- Both the directory structure and the files reside on disk.
- Backups of these two structures are kept on tapes.

A Typical File-system Organization



Information in a Device Directory

- Name
- Type
- Address
- Current length

- Maximum length
- Date last accessed (for archival)
- Date last updated (for dump)
- Owner ID (who pays)
- Protection information (discuss later)

Operations Performed on Directory


- Search for a file
- Create a file
- Delete a file
- List a directory
- Rename a file
- Traverse the file system

Organize the Directory (Logically) to Obtain

- **Efficiency** (效率) – locating a file quickly.
- **Naming** (命名) – convenient to users.
 - ☞ Two users can have same name for different files.
 - ☞ The same file can have several different names.
- **Grouping** (组) – logical grouping of files by properties, (e.g., all Java programs, all games, ...)

Single-Level Directory


- A single directory for all users.



Naming problem
Grouping problem


Two-Level Directory

- Separate directory for each user.



- Path name（路径名）
- Can have the same file name for different user
- Efficient searching
- No grouping capability

Tree-Structured Directories (树形结构目录)



Tree-Structured Directories (Cont.)

- Efficient searching
- Grouping Capability
- Current directory (working directory)
 - ➡ **cd** /spell/mail/prog
 - ➡ **type** list

- **Absolute or relative** path name
- Creating a new file is done in current directory.
- Delete a file
rm <file-name>
- Creating a new subdirectory is done in current directory.
mkdir <dir-name>

Tree-Structured Directories (Cont.)

Example: if in current directory **/mail**


mkdir count



Deleting “mail” $\Rightarrow$ deleting the entire subtree rooted by “mail”.

(无环图目录)

- Have shared subdirectories and files.




Acyclic-Graph Directories (Cont.)

- Two different names (aliasing)
- If *dict* deletes *list* $\Rightarrow$ dangling pointer.

Solutions:

- ☞ Backpointers, so we can delete all pointers.
Variable size records a problem.
- ☞ Backpointers using a daisy chain organization.
- ☞ Entry-hold-count solution.

General Graph Directory (通用图目录)




General Graph Directory (Cont.)

- How do we guarantee no cycles?
 - ☞ Allow only links to file not subdirectories.
 - ☞ Garbage collection.
 - ☞ Every time a new link is added use a cycle detection algorithm to determine whether it is OK.


File System Mounting

- A file system must be **mounted** before it can be accessed.
- A unmounted file system (l.e. Fig. 11-11(b)) is mounted at a **mount point**.

(a) Existing. (b) Unmounted Partition




(a)



(b)

Mount Point (安装点)



File Sharing

- Sharing of files on multi-user systems is desirable.
- Sharing may be done through a *protection* scheme.
- On distributed systems, files may be shared across a network.
- Network File System (NFS) is a common distributed file-sharing method.

Protection

- File owner/creator should be able to control:
 - ☞ what can be done
 - ☞ by whom

■ Types of access

➡ Read

➡ Write

➡ Execute

➡ Append

➡ Delete

➡ List

- Mode of access: read, write, execute
- Three classes of users

RWX

a) **owner access** 7 $\Rightarrow$ 1 1 1

RWX

b) **group access** 6 $\Rightarrow$ 1 1 0

RWX

c) **public access** 1 $\Rightarrow$ 0 0 1

Access Lists and Groups

(访问列表和组)

- Ask manager to create a group (unique name), say G, and add some users to the group.
- For a particular file (say *game*) or subdirectory, define an appropriate access.

owner group public

chmod 761 game

Attach a group to a file

chgrp G game